

CoachAI — File Structure (Short)

Main file: run.py | **UI file:** src/ui/app.py | **Prompts:** src/core/prompts.py | **Scores:** data/reports/

src/core/ — control room

File	Job
orchestrator.py	Starts/ends session, routes each turn
llm.py	Calls Groq LLM (with fallback models)
prompts.py	Every agent's system prompt, in one file
models.py / config.py / database.py	Data schemas / settings / SQLite save

src/modules/ — per-turn logic

File	Job
conversation_manager.py	Runs agents in order each turn
session_config.py	Simulator / Manual / Replay mode setup
performance_analytics.py	Score trends across past sessions

src/agents/ — 24 AI agents (one job each)

File	Job
customer_simulator.py	Generates the fake customer
intent_sentiment.py	Reads mood/intent/frustration
knowledge_recommendation.py	RAG — finds KB answer
coaching_suggestion.py	Suggests what agent should say
escalation_monitor.py	Scores escalation risk
post_interaction_summary.py	Builds the final score/report
+ 18 more	Bonus agents: CSAT, fraud, compliance, tone, etc.

src/rag/ and src/ui/

File	Job
rag/knowledge_base.py	Chunks files + keyword search
ui/app.py	Main Streamlit UI — all tabs
ui/panels.py	Chat + coaching panel rendering

data/ — where scores/content live

File	Job
data/reports/	Score/report JSON per session
data/coach.db	Session database
data/knowledge_base/	FAQ docs the RAG searches

Say this: "Main file is run.py, UI is in src/ui/app.py, all AI logic is in 24 files under src/agents/, src/core/ handles orchestration + prompts, scores save to data/reports/."